

CARA OUELLETTE

Toronto, ON | 613-447-3930 | cara.ouellette@georgebrown.ca | github.com/CaraOuellette | timetopivot.substack.com

OBJECTIVE

Full-Stack ML Engineer with 4+ years in enterprise implementation and production support, plus hands-on applied AI development. Built a prototype agent harness for Cohere's Command A LLM with tool/function calling, RAG memory, and intensive prompt iteration. Strong coding background across Python, SQL, Docker, and web stacks. Currently completing a postgraduate program in Applied AI Solutions Development (George Brown Polytechnic, Dec 2026). Seeking a Prompt Engineer internship focused on model persona in AI — GTA or National Capital Region, September 2026.

HIGHLIGHT OF QUALIFICATIONS

- Comprehensive familiarity with all Cohere models, including adversarial/red-team techniques
 - Recognized contributor in cyborgist, rationalist, and AI circles; Hyperplex SF residency participant (2025)
 - Prodigious with Claude Code/Codex; experienced building personal agent harnesses and RAG systems
 - Completed and published multiple independent research projects on AI persona and LLM behavior
 - Serious, curious, and fully devoted to pushing the frontier of AI knowledge
-

TECHNICAL SKILLS

AI Competencies: Prompt Engineering, Model Evaluation, Fine-tuning, Red-teaming

Programming: Python, SQL, JavaScript, PHP (Laravel), Git, Docker

Data/Infra: Postgres, MySQL/MariaDB, MongoDB, Visual Studio

Business Ops: Azure DevOps, Jira, Zendesk, Confluence

EDUCATION

Postgraduate Certificate (In Progress) — Applied AI Solutions Development, George Brown Polytechnic | Expected Dec 2026

Focus: model training & fine-tuning, applied ML math, big data workflows, production AI solution development

BSc (Honours) — Mathematics & Statistics, University of Toronto | 2020

EXPERIENCE

Timeless Medical Systems — Charlottetown, PE | May 2021 – Jun 2025

Support Technician / Implementation Engineer | Jan 2023 – Jun 2025 - Delivered implementation and production support for hospital software in PHP/Laravel + SQL environments; diagnosed complex incidents, verified fixes, maintained high-quality escalation notes and RCAs. - Completed ~6 overlapping multi-month implementations with 100% deployment success, coordinating interface testing and stakeholder validation. - Wrote code daily: client-specific EMR processing components, SQL scripts, and diagnostic utilities to accelerate triage. - Produced runbooks and support-to-implementation documentation that improved handoffs and reduced time-to-triage for recurring issues.

Support Analyst | May 2021 – Jan 2023 - Operated in high-backlog conditions (~100 active cases), sustaining SLA-driven triage and clear partner communication. - Recognized as consistent top performer (top 10% internal ranking).

Independent AI Researcher / Journalist | Jun 2025 – Present - Publishes accessible analysis of AI systems and agentic workflows; building toward a longer technical writeup of ongoing LLM research.

SELECTED PROJECTS

Command A Agent Harness — Prompt & Agent Engineering - Built an agent harness for well-defined programming tasks requiring tool use and local filesystem manipulation (planning → execution → verification → output). - Implemented tool/function calling with structured outputs, schema constraints, retries, and validation. - Developed a RAG-style memory system (Postgres + pgvector) storing conversation summaries, key facts, preferences, and tool outputs; retrieval used semantic similarity + recency weighting.

LLM Research Project — Positional Bias Reduction (Prompting + Evaluation) - Designed a controlled experiment to reduce positional bias using counterbalanced options and rubric-based scoring. - Built a lightweight evaluation harness to compare prompt variants and document failure modes and tradeoffs. - Achieved ~20% reduction in positional preference errors on a synthetic evaluation set.